

Report of Practical Course on High-Performance Computing

Parallel Deep Learning pipelines using Go and MPI

Persentor: Silin Zhao Supervisor: Patrick Michaelis

September 27, 2022

Project notation

Youtube link

<https://www.youtube.com/watch?v=2siZQBvRPuY&t=6s>

Datasets

- ▶ This project source code can be found https://github.com/scofild429/go_mpi_network, This is the **README** page.
- ▶ Iris dataset (<https://www.kaggle.com/datasets/saurabh00007/iriscsv>)
- ▶ Intel image classification, (<https://www.kaggle.com/datasets/puneet6060/intel-image-classification?resource=download>). Download it, put archive it in the folder `./datasets/`

All training data will equally divided for each training network, specially for mpi

Configuration example

- ▶ `./goai/.irisenv`
- ▶ `./goai/.imgenv`

```
inputdataDims=4
inputLayerNeurons=30
hiddenLayerNeurons=20
outputLayerNeurons=3
labelOnehotDims=3
numEpochs=100
learningRate=0.01
batchSize=4
```

Submit the job in cluster

no singularity, installing golang 1.18 was failed always using binary executable code of golang, **go build**

```
#!/bin/bash
#SBATCH --job-name mpi-go-neural-network
#SBATCH -N 1
#SBATCH -p fat
#SBATCH -n 20
#SBATCH --time=01:30:00

module purge
module load openmpi

mpirun -n 20 ./goai
```

Deep learning's problem

As AI comes to deep learning, the computing resource becomes more critical for training process.

Applications:

- ▶ Image Classification
- ▶ NLP
- ▶ Semantic segmentation

Solution

- ▶ GPU
- ▶ TPU
- ▶ **Distributed learning**

Single network architecture

raining data \rightarrow inputLayer(w1, b1) \rightarrow dinputLayer

Normalization

dinputLayer \rightarrow hiddenLayer(w2, b2) \rightarrow dhiddenLayer

Normalization

dhiddenLayer \rightarrow OutputLayer(w3, b3) \rightarrow doutputLayer

Loss = L2: $(\text{doutputLayer} - \text{onehotlable})^2$

Backpropagation from Loss of Outputlayer to w3, b3

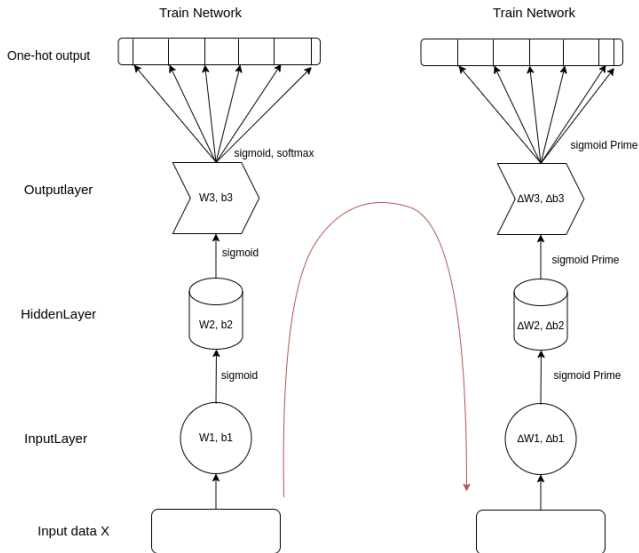
Backpropagation from error of Hiddenlayer to w2, b2

Backpropagation from error of Inputlayer to w1, b1

Derivative of sigmoid, Normalization, Standardization

- ▶ Stochastic Gradient Descent (SGD)
- ▶ Mini-batch Gradient Descent (MBGD)
- ▶ Batch Gradient Descent (BGD)

Illustration of weights updating



Code implementation

```
func main() {  
  ^^Isinglenode.Single_node_iris(true)  
  ^^Impicode.Mpi_iris_Allreduce()  
  ^^Impicode.Mpi_iris_SendRecv()  
  ^^Impicode.Mpi_images_Allreduce()  
  ^^Impicode.Mpi_images_SendRecv()  
}
```


You can review my code, and choose one of them to be executed in `/goai/myai.go` main function.

Comparing with python:


- ▶ `./pytorchDemo/irisfromscratch.py`
- ▶ `./pytorchDemo/iriswithpytorch.py`
- ▶ `./pytorchDemo/logisticRcuda.py`

Network performance(iris dataset)

Loss



Accuracy



MPI communication

```
github.com/sbromberger/gompi  
import CGO as C
```


► **Collective**

- `gompi.BcastFloat64s()` -> `C.MPI_Bcast()`
- `gompi.AllreduceFloat64s` -> `C.MPI_Allreduce()`

► **Non Collective**

- `gompi.SendFloat64s()` -> `C.MPI_Send()`
- `gompi.SendFloat64()` -> `C.MPI_Send()`
- `gompi.RecvFloat64s()` -> `C.MPI_Recv()`
- `gompi.RecvFloat64()` -> `C.MPI_Recv()`

Non collective architecture



Non collective design

rank = 0

- ▶ in **main network** weights will be initialized, but not for training,
- ▶ weights will broadcast to all other training networks

rank != 0

- ▶ in **train network** receive weights from main network for initialization
- ▶ After each batch training done, sending its weights variance to main network

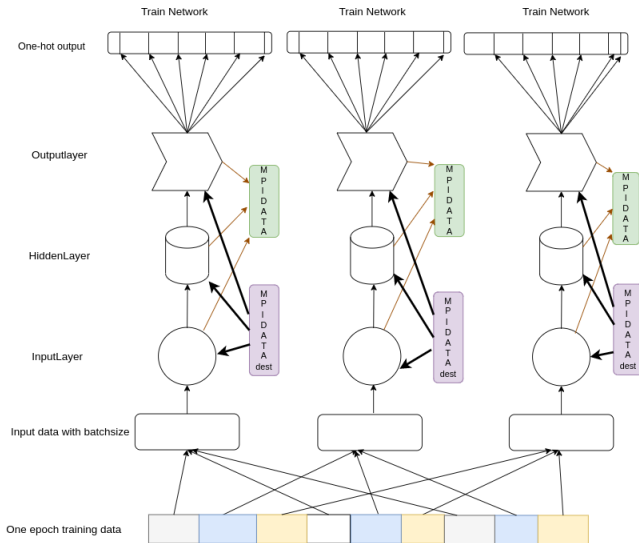
rank = 0

- ▶ receiving the variance from all training network
- ▶ accumulating and then sending back to training network

rank != 0

- ▶ start next training batch

Collective architecture




Collective design


- ▶ All network train its data respectively,
- ▶ After each train batch, pack all weights into array
- ▶ $\text{MPI}_{\text{Allreduce}}$ for new array
- ▶ updating weights with new array

Iris dataset performance for non-collective

Send&Recv loss




Send&Recv accuracy




Iris dataset performance for collective

Allreduce loss




Allreduce accuracy




Intel image classification performance


Send&Recv loss (220 images)




Allreduce loss (220 images)



SendRecv loss (14000 images)




Allreduce loss (14000 images)




Speedup Diagrams

Iris for Allreduce and Send&Recv with different nodes



Intel Image Classification for Allreduce and Send&Recv with different nodes



Discussion

neural network model implement is not perfect, so the accuracy performance not so well

For each epoch:

- ▶ Allreduce: about 2 minutes
- ▶ Send&Recv: about 3.6 minutes, because of synchronization of each batch training

Change nodes, scaling behavior, such as speedup diagrams is missing

Change the batchsize, reducing mpi communication

Conclusion

- ▶ Golang can also be used for parallel computing
- ▶ neural network implementation of golang can be improved
- ▶ HPC cluster for distributed learning has significant benefits for large dataset